

CSCE 580: Introduction to AI *CSCE 581: Trusted AI*

Lecture 6: Search

PROF. BIPLAV SRIVASTAVA, AI INSTITUTE

12TH SEP 2023

Carolinian Creed: “I will practice personal and academic integrity.”

Credits: Copyrights of all material reused acknowledged

Organization of Lecture 6

- Introduction Segment
 - Recap of Lecture 5
- Main Segment
 - Problem solving agent – goal directed
 - Problem formulation – abstraction, type of problems
 - Search
- Concluding Segment
 - Course Project Discussion
 - About Next Lecture – Lecture 7
 - Ask me anything

Sprint 1: (Sep 12 – Oct 5) STARTS

Introduction Section

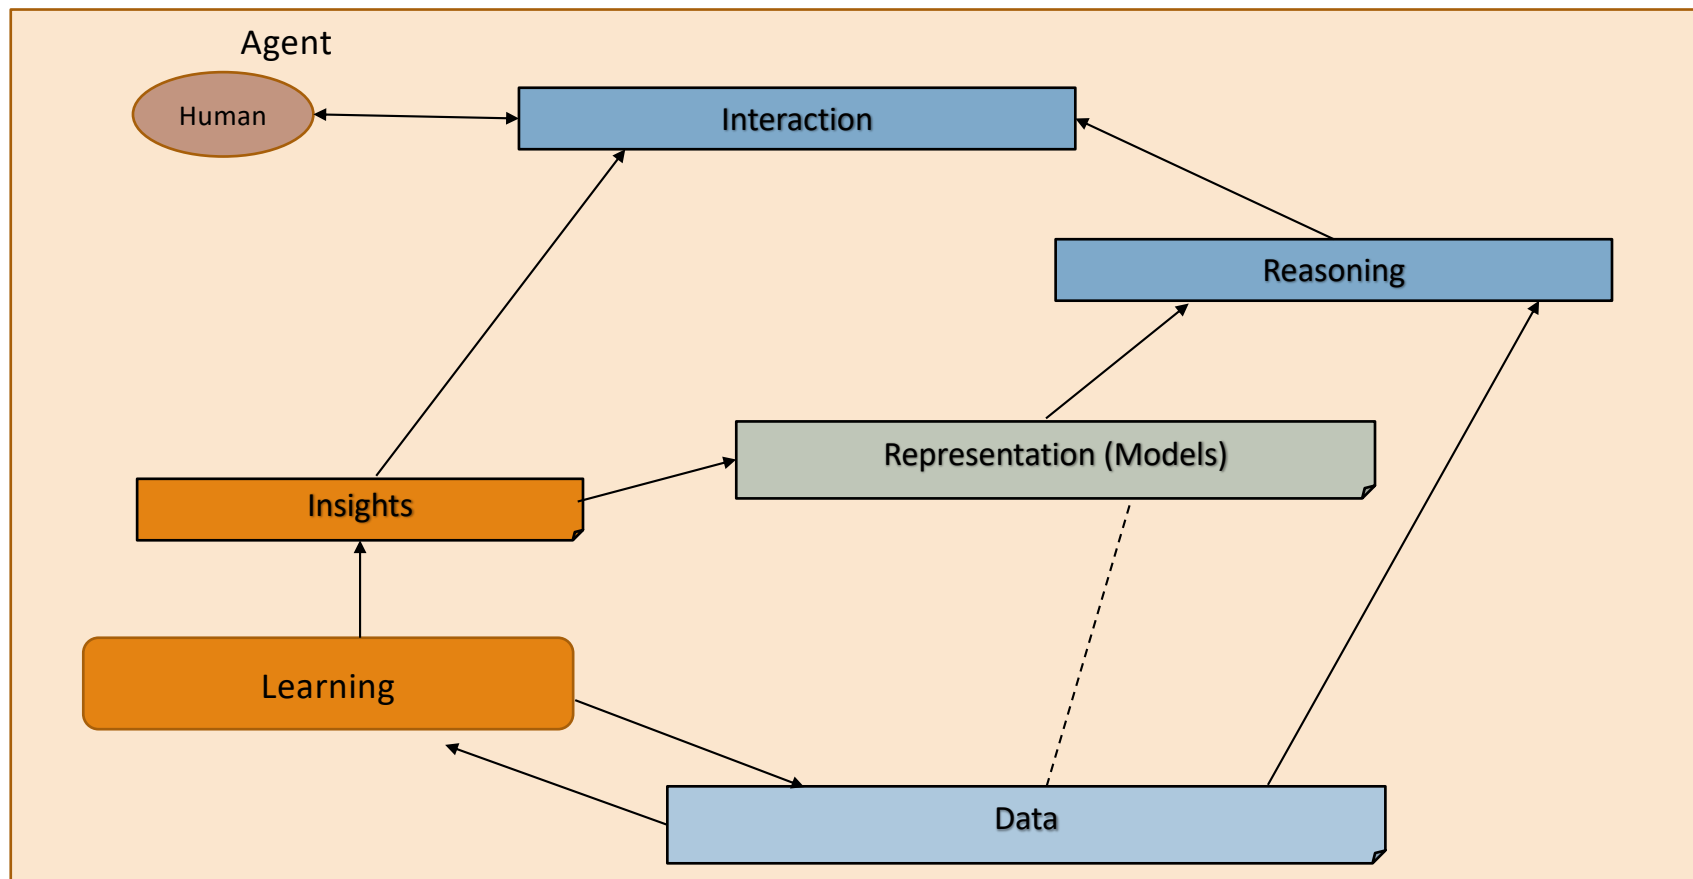
Recap of Lecture 5

- Logic – First Order
- Inferencing
- Representation in the Large: ConceptNet, Cyc
- Trust Issues with Knowledge Representation

Intelligent Agent Model



Relationship Between Main AI Topics



Where We Are in the Course


CSCE 580/ 581 – In This Course

- Week 1: Introduction, Aim: Chatbot / Intelligence Agent
- Weeks 2-3: Data: Formats, Representation and the Trust Problem
- Week 4-5: Search, Heuristics - Decision Making
- Week 6: Constraints, Optimization – Decision Making
- Week 7: Classical Machine Learning – Decision Making, Explanation
- Week 8: Machine Learning - Classification
- Week 9: Machine Learning - Classification – Trust Issues and Mitigation Methods
- Topic 10: Learning neural network, deep learning, Adversarial attacks
- Week 11: Large Language Models – Representation, Issues
- Topic 12: Markov Decision Processes, Hidden Markov models - Decision making
- Topic 13: Planning, Reinforcement Learning – Sequential decision making
- Week 14: AI for Real World: Tools, Emerging Standards and Laws; Safe AI/ Chatbots

Main Section

Example: Vacuum World

- Situation
 - Two rooms
 - One robot
 - Dirt can be in any room
- Goal
 - Clean the rooms
- Actions
 - Move left, move right, clean

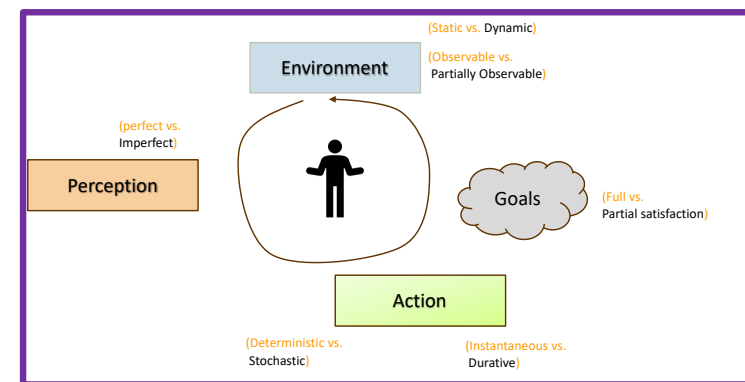


Adapted from:

1. Russell & Norvig, AI: A Modern Approach
2. Bart Selman's CS 4700 Course

Goal-directed Problem Solving Agents

1. Goal Formulation: Have one or more (desirable) world states
2. Problem formulation: What actions and states to consider given goals and an initial state
3. Search for solution: Given the problem, search for a solution - a sequence of actions to achieve the goal starting from the initial state
4. Execution: agent can execute actions in the solution




Adapted from:

1. Russell & Norvig, AI: A Modern Approach
2. Bart Selman's CS 4700 Course

Modeling and Abstraction Consideration

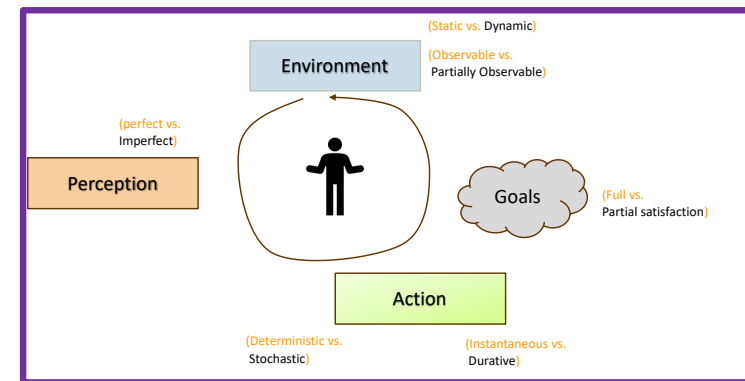
- Model: an abstract representation of the problem
 - “All models are wrong, but some are useful”
- What to capture, what to avoid
 - Only the necessary details needed to solve the problem
- In the example, we can avoid
 - For concepts
 - Size of rooms or robot
 - Quantity of dirt
 - For actions
 - Time taken to clean
 - Charging/ recharging time
 - Doing nothing – staying at the same place?



- Concepts
 - Two rooms
 - One robot
 - Dirt can be in any room
- Goal
 - Clean the rooms
- Actions
 - Move left, move right, clean

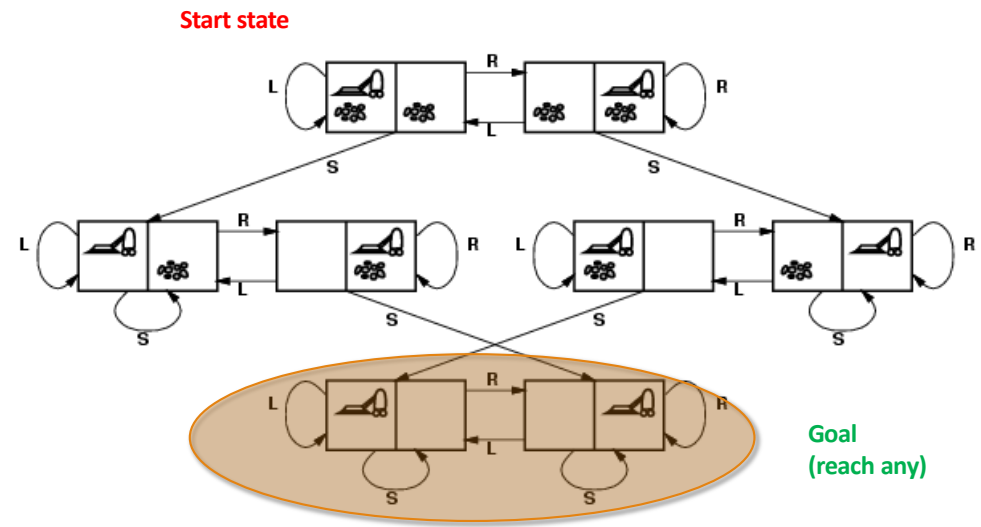
Open Loop v/s Closed Loop Systems

- Open loop
 - Assuming the world will not change, after a solution is found, one can simply execute it one action at a time
- Closed loop
 - If the world keeps changing, after a solution is found, one cannot ignore perception when executing actions
 - The solution has to be relooked whenever an action is being executed. New solutions may have to be found at each step again.



Formulating a Problem

States	8 possible world states <i>(2room x 2dirt location x 2clean?)</i>
<ul style="list-style-type: none"> Initial state Goal state 	<ul style="list-style-type: none"> Any No dirt at all locations
Actions	Left, Right, Suck
<ul style="list-style-type: none"> Transition model Action cost 	<ul style="list-style-type: none"> Action transition (edges) 1



Adapted from:

1. Russell & Norvig, AI: A Modern Approach
2. Bart Selman's CS 4700 Course

Type of Problems

Standardized Problems

- Grid world
- Sliding tile
- Sokoban
- Chess
- ...

Real-World Problems

- Route finding
- Robotic / space craft navigation
- Protein design: find a sequence of amino acids that will fold into a 3D protein structure
- Dialog generation: how to give an effective answer that a person can understand
- ...

Exercise: Sliding 8-tile Puzzle

States

- Initial state
- Goal state

Actions

- Transition model
- Action cost

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Adapted from:
Russell & Norvig, AI: A Modern Approach

Exercise: Sliding 8-tile Puzzle

States	Location of tiles
<ul style="list-style-type: none">Initial stateGoal state	<ul style="list-style-type: none">Any (given)All numbers sorted, Empty tile in corner (given)
Actions	move blank left, right, up, down
<ul style="list-style-type: none">Transition modelAction cost	<ul style="list-style-type: none">Blank transition (edges)1

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State


Adapted from:

1. Russell & Norvig, AI: A Modern Approach
2. Bart Selman's CS 4700 Course

Search

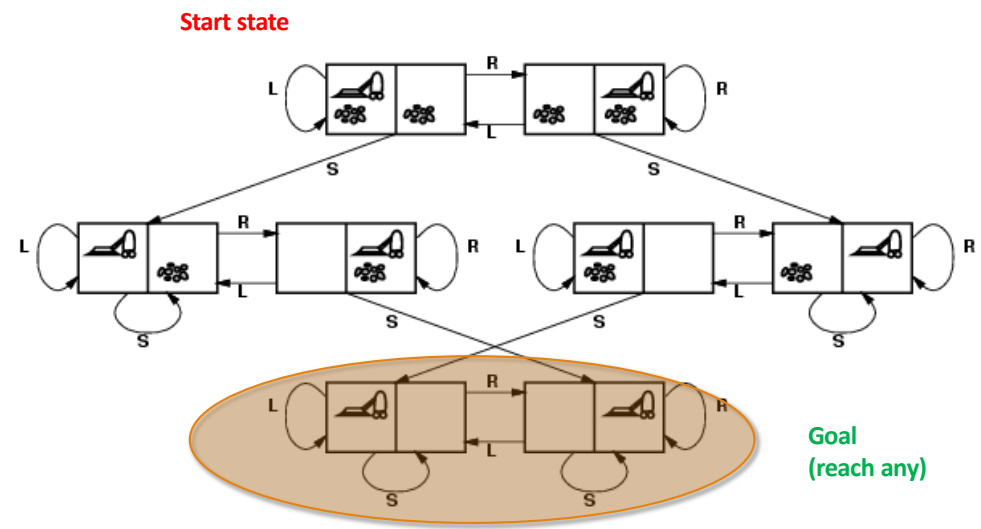

Search Basics

- input: a problem with states and actions
- output: solution(s) or flag for failure
- Concepts:
 - **Node**: corresponds to a state of the problem
 - **Edges**: transition between states
 - **Expand**: consider actions in the state (ACTIONS) and transition model. Generate new nodes corresponding to resulting states (RESULT)
 - **Explore**: check when a node meets goal condition



Node 1 has been **Reached**, Nodes {2,3 4} constitute Node 1's **Frontier**

Formulating a Problem



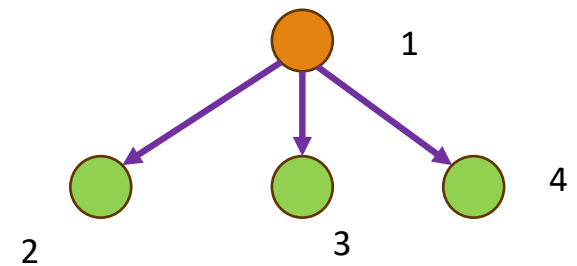
Adapted from:

1. Russell & Norvig, AI: A Modern Approach
2. Bart Selman's CS 4700 Course

Tree-search Algorithms

Basic idea: simulated exploration of state space by generating successors of already-explored states (a.k.a. ~ **expanding** states)

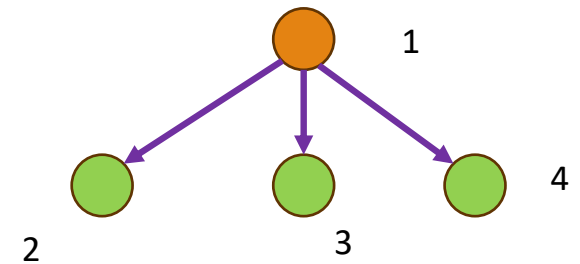
```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```



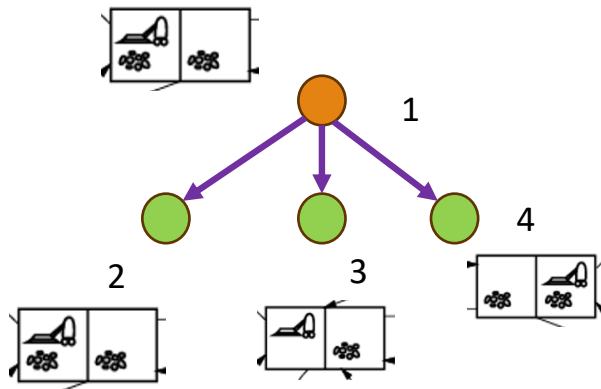
Adapted from:
1. Russell & Norvig, AI: A Modern Approach
2. Bart Selman's CS 4700 Course

Implementing a Tree-Search Algorithm

- Data structure
 - node.STATE: the state to which the node corresponds
 - node.PARENT: the node in the tree that generated this node
 - node.ACTION: the action that was applied to the parent to generate this node
 - node.PATH-COST: the total cost of the path from the initial state to this node
- Queue to store frontier
 - Is-EMPTY(frontier): true/ false depending on whether frontier is empty
 - POP(frontier): remove the top node from the frontier and return it
 - TOP(frontier): returns the top node from the frontier but does not remove it
 - ADD(node, frontier): insert node into its proper place in the queue
- Queue:
 - priority queue – removes the node with minimum cost according to some evaluation function
 - FIFO queue – first in, first output. Used in breadth first search
 - LIFO queue - last in, first output. Used in depth first search



Best-First Search



```

function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure
  node  $\leftarrow$  NODE(STATE=problem.INITIAL)
  frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
  reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    if problem.IS-GOAL(node.STATE) then return node
    for each child in EXPAND(problem, node) do
      s  $\leftarrow$  child.STATE
      if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
        reached[s]  $\leftarrow$  child
        add child to frontier
  return failure
  
```

```

function EXPAND(problem, node) yields nodes
  s  $\leftarrow$  node.STATE
  for each action in problem.ACTIONS(s) do
    s'  $\leftarrow$  problem.RESULT(s, action)
    cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
    yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
  
```

Examples of Search Strategies

- Uninformed
 - Depth first
 - Breadth first
- Informed (Heuristic)
 - Greedy best first search
 - A* search

More on Search Strategies

- A search strategy is defined by picking the **order of node expansion**.
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - *b*: maximum branching factor of the search tree
 - *d*: depth of the least-cost solution
 - *m*: maximum depth of the state space (may be ∞)

Adapted from:

1. Russell & Norvig, AI: A Modern Approach
2. Bart Selman's CS 4700 Course

Exercise and Code

- Search Methods
 - From Book: AI – A Modern Approach,
<https://github.com/aimacode/aima-python/blob/master/search.ipynb>

Source: Russell & Norvig, AI: A Modern Approach

Lecture 6: Summary

- We talked about
 - Goal-directed problem solving agents
 - How to formulate problem formulations
 - Search concepts
 - Problems of controlled robot navigation, 8-tile
 - Search strategies

Concluding Section

Course Project

Project Discussion: What Problem Fascinates You ?

- Data
 - Water
 - Finance
 - ...
- Analytics
 - Search, Optimization, Learning, Planning, ...
- Application
 - Building chatbot
- Users
 - Diverse demographics
 - Diverse abilities
 - Multiple human languages

Project execution in sprints

- Sprint 1: (Sep 12 – Oct 5)
 - **Solving**: Choose a decision problem, identify data, work on solution methods
 - **Human interaction**: Develop a basic chatbot (no AI), no problem focus
- Sprint 2: (Oct 10 – Nov 9)
 - **Solving**: Evaluate your solution on problem
 - **Human interaction**: Integrated your choice of chatbot (rule-based or learning-based) and methods
- Sprint 3: (Nov 14 – 30)
 - **Evaluation**: Comparison of your solver chatbot with an LLM-based alternative, like ChatGPT

Project Discussion: Dates and Deliverables

Project execution in sprints

- Sprint 1: (Sep 12 – Oct 5)
 - **Solving**: Choose a decision problem, identify data, work on solution methods
 - **Human interaction**: Develop a basic chatbot (no AI), no problem focus
- Sprint 2: (Oct 10 – Nov 9)
 - **Solving**: Evaluate your solution on problem
 - **Human interaction**: Integrated your choice of chatbot (rule-based or learning-based) and methods
- Sprint 3: (Nov 14 – 30)
 - **Evaluation**: Comparison of your solver chatbot with an LLM-based alternative, like ChatGPT

- Oct 12, 2023
 - Project checkpoint
 - In-class presentation
- Nov 30, 2023
 - Project report due
- Dec 5 / 7, 2023
 - In-class presentation

Skeleton: A Basic Chatbot

- Run in an infinite loop until the user wants to quit
- Handle any user response
 - User can quit by typing “Quit” or “quit” or just “q”
 - User can enter any other text and the program has to handle it. The program should write back what the user entered and say – “I do not know this information”.
- Handle known user query types // Depends on your project
 - “Tell me about N-queens”, “What is N ?”
 - “Solve for N=4?”
 - “Why is this a solution? ”
- Handle chitchat // Support at least 5, extensible from a file
 - “Hi” => “Hello”
 - ...
- *Store session details in a file*

Illustrative Project

1. **Title:** Solve and explain solving of n-queens puzzle
2. **Key idea:** Show students how a course project will look like
3. **Who will care when done:** students of the course, prospective AI students and teachers
4. **Data need:** n: the size of game; interaction
5. **Methods:** search
6. **Evaluation:** correctness of solution, quality of explanation, appropriateness of chat
7. **Users:** with and without AI background; with and without chess background
8. **Trust issue:** user may not believe in the solution, may find interaction offensive (why queens, not kings? ...)

Project Discussion: Illustration

1. Create a private Github repository called “CSCE58x-Fall2023-<studentname>-Repo”. Share with Instructor (biplav-s) and TA (kausik-l)
2. Create Google folder called “CSCE58x-Fall2023-<studentname>-SharedInfo”. Share with Instructor (prof.biplav@gmail.com) and TA (lakkarajukausk90@gmail.com)
3. Create a Google doc in your Google repo called “Project Plan” and have the following by next class (Sep 5, 2023)

1. **Title:** Solve and explain solving of n-queens puzzle
2. **Key idea:** Show students how a course project will look like
3. **Who will care when done:** students of the course, prospective AI students and teachers
4. **Data need:** n: the size of game; interaction
5. **Methods:** search
6. **Evaluation:** correctness of solution, quality of explanation, appropriateness of chat
7. **Users:** with and without AI background; with and without chess background
8. **Trust issue:** user may not believe in the solution, may find interaction offensive (why queens, not kings? ...)

Reference: Project Rubric

- **Project results – 60%**
 - Working system ? – 30%
 - Evaluation with results superior to baseline? – 20%
 - Considered related work? – 10%
- **Project efforts – 40%**
 - Project report – 20%
 - Project presentation (updates, final) – 20%
- **Bonus**
 - Challenge level of problem – 10%
 - Instructor discretion – 10%
- **Penalty**
 - Lack of timeliness as per announced policy (right) - up to 30%

Milestones and Penalties

- Oct 12, 2023
 - Project checkpoint
 - In-class presentation
 - **Penalty: presentation not ready by Oct 10, 2023 [-10%]**
- Nov 30, 2023
 - Project report due
 - **Project report not ready by date [-10%]**
- Dec 5 / 7, 2023
 - In-class presentation
 - **Project presentations not ready by Dec 4, 2023 [-10%]**

About Next Lecture – Lecture 7

Lecture 7: Searching for Problem Solving

- Search
- Heuristics