

CSCE 580: Introduction to AI *CSCE 581: Trusted AI*

Lecture 10: Game Search

PROF. BIPLAV SRIVASTAVA, AI INSTITUTE

26TH SEP 2023

Carolinian Creed: “I will practice personal and academic integrity.”

Credits: Copyrights of all material reused acknowledged

Organization of Lecture 10

- Introduction Segment
 - Recap of Lecture 9
- Main Segment
 - Games and Search
 - Minimax
 - Alphabeta
 - Monte Carlo Tree Search
- Concluding Segment
 - Course Project Discussion
 - About Next Lecture – Lecture 11
 - Ask me anything

Introduction Section

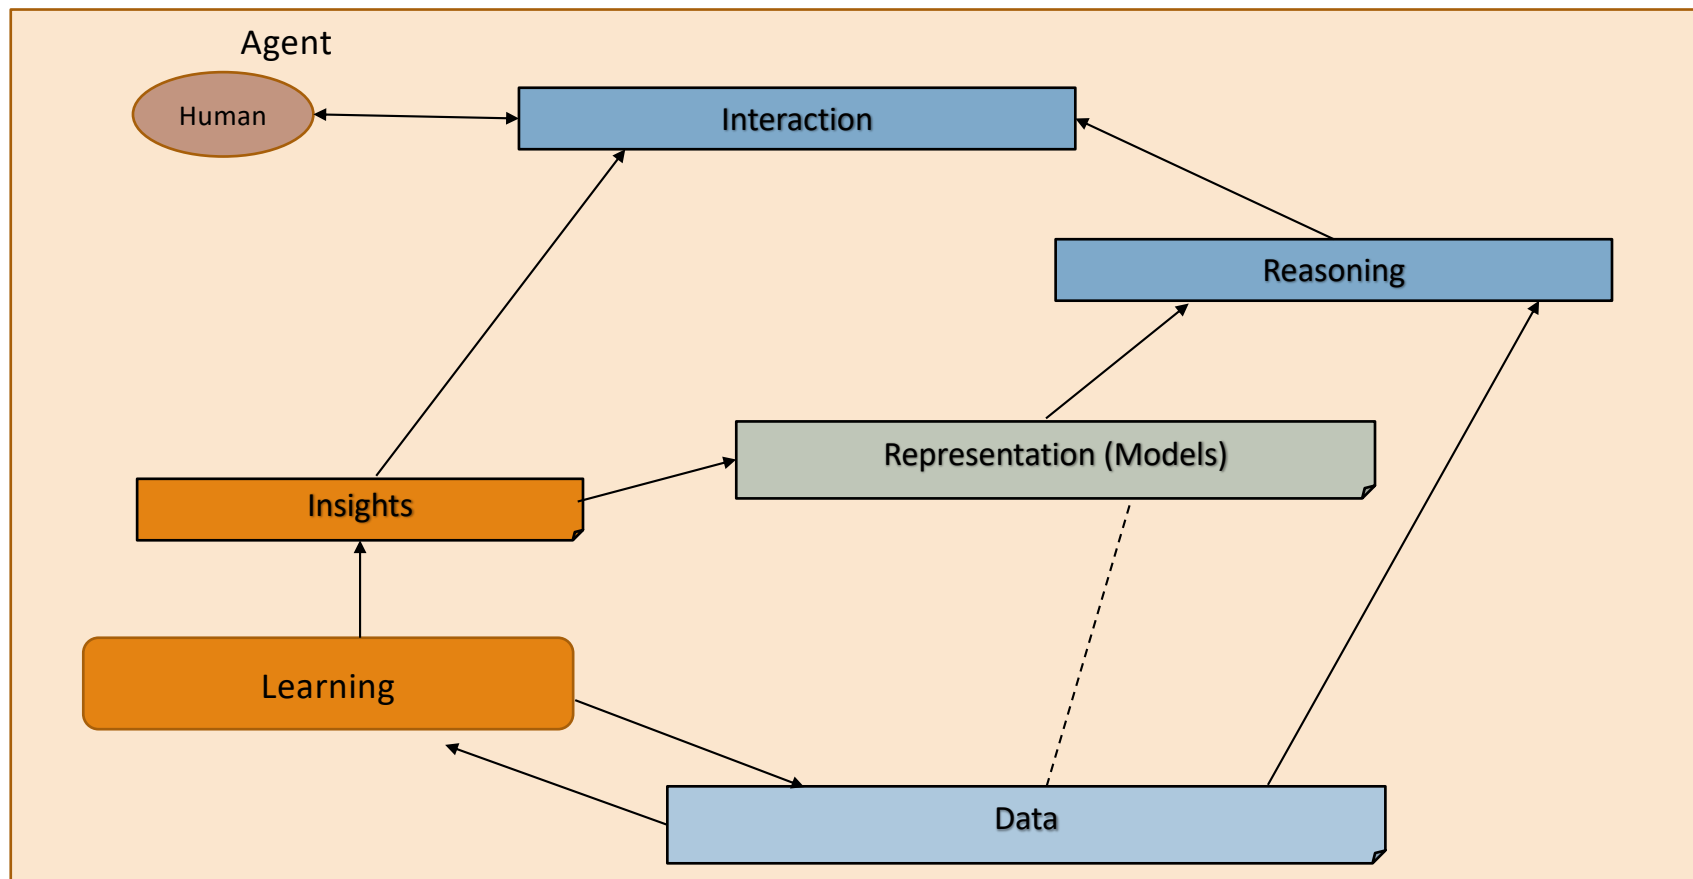
Recap of Lecture 10

- Hill climbing
- Simulated Annealing
- Genetic programming
- Search in complex environments

Intelligent Agent Model



Relationship Between Main AI Topics



Where We Are in the Course

CSCE 580/ 581 – In This Course

- Week 1: Introduction, Aim: Chatbot / Intelligence Agent
- Weeks 2-3: Data: Formats, Representation and the Trust Problem
- Week 4-5: Search, Heuristics - Decision Making
- Week 6: Constraints, Optimization – Decision Making
- Week 7: Classical Machine Learning – Decision Making, Explanation
- Week 8: Machine Learning - Classification
- Week 9: Machine Learning - Classification – Trust Issues and Mitigation Methods
- Topic 10: Learning neural network, deep learning, Adversarial attacks
- Week 11: Large Language Models – Representation, Issues
- Topic 12: Markov Decision Processes, Hidden Markov models - Decision making
- Topic 13: Planning, Reinforcement Learning – Sequential decision making
- Week 14: AI for Real World: Tools, Emerging Standards and Laws; Safe AI/ Chatbots

Main Section

Offline and Online Search

- All search studied until now were offline search
 - Solution found and then agent executed
- Online search
 - Interleave solution finding and execution
 - Cannot guarantee solution, unless actions have inverse-actions to undo state change


Games and Search

- Common setting
 - Two players
 - One human, one automated [Common case]
 - Both automated
 - Perfect information (fully observable)
 - Zero-sum – if one wins, another loses
- Captures many types of games
 - Tic-tac-toe, chess, Go, backgammon, ...

Games and Search

- Setup
 - S_0 : the initial state
 - $TO-MOVE(s)$: the player to play in state s
 - $ACTIONS(s)$: the set of legal actions at state s
 - $RESULT(s, a)$: transition model
 - $IS-TERMINAL(s)$: terminal test to determine if game is over
 - $UTILITY(s, p)$: utility or objective function. Numeric value capturing value of state s to player p

Tic-Tac-Toe – Game Tree



Adapted from:
Russell & Norvig, AI: A Modern Approach


MiniMax Utility

MINMAX(s) =

- $UTILITY(s, MAX)$ if IS-TERMINAL(s)
- $\text{Max } a \text{ in } ACTIONS(s) \text{ MINIMAX}(RESULT(s, a))$ if TO-MOVE(s) = MAX
- $\text{Min } a \text{ in } ACTIONS(s) \text{ MINIMAX}(RESULT(s, a))$ if TO-MOVE(s) = MIN

MAX

MIN



Adapted from:
Russell & Norvig, AI: A Modern Approach

Minimax Search Algorithm

```
function MINIMAX-SEARCH(game, state) returns an action
  player  $\leftarrow$  game.TO-MOVE(state)
  value, move  $\leftarrow$  MAX-VALUE(game, state)
  return move

function MAX-VALUE(game, state) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v  $\leftarrow -\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MIN-VALUE(game, game.RESULT(state, a))
    if v2 > v then
      v, move  $\leftarrow$  v2, a
  return v, move

function MIN-VALUE(game, state) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v  $\leftarrow +\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MAX-VALUE(game, game.RESULT(state, a))
    if v2 < v then
      v, move  $\leftarrow$  v2, a
  return v, move
```


Starts from a Max node
Recursively does min on children,
who in-turn does max on children
to get value and corresp. action

Adapted from:
Russell & Norvig, AI: A Modern Approach

Multi-player (>2) Games

- Possible to extend MINMAX, algorithm is more complex
 - Extend values with vectors, corresponding to per player
 - Levels increase (per turn)
- In real games, players can often form alliances
 - Hard to model


Alpha-Beta Pruning



If m or m' is better than n for Player,
search will never get to n

Adapted from:
Russell & Norvig, AI: A Modern Approach

Alpha-Beta Pruning



The first leaf below C has the value 2. Hence, C, which is a MIN node, has a value of at most 2. But we know that B is worth 3, so MAX would never choose C.

Adapted from:
Russell & Norvig, AI: A Modern Approach

Alpha-Beta Pruning

function ALPHA-BETA-SEARCH(*game*, *state*) **returns** an action

player \leftarrow *game*.TO-MOVE(*state*)

value, move \leftarrow MAX-VALUE(*game*, *state*, $-\infty$, $+\infty$)

return move

function MAX-VALUE(*game*, *state*, α , β) **returns** a (utility, move) pair

if *game*.IS-TERMINAL(*state*) **then return** *game*.UTILITY(*state*, player), null

$v \leftarrow -\infty$

for each *a* **in** *game*.ACTIONS(*state*) **do**

$v2, a2 \leftarrow$ MIN-VALUE(*game*, *game*.RESULT(*state*, *a*), α , β)

if $v2 > v$ **then**

$v, move \leftarrow v2, a$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

if $v \geq \beta$ **then return** $v, move$

return $v, move$

function MIN-VALUE(*game*, *state*, α , β) **returns** a (utility, move) pair

if *game*.IS-TERMINAL(*state*) **then return** *game*.UTILITY(*state*, player), null

$v \leftarrow +\infty$

for each *a* **in** *game*.ACTIONS(*state*) **do**

$v2, a2 \leftarrow$ MAX-VALUE(*game*, *game*.RESULT(*state*, *a*), α , β)

if $v2 < v$ **then**

$v, move \leftarrow v2, a$

$\beta \leftarrow \text{MIN}(\beta, v)$

if $v \leq \alpha$ **then return** $v, move$

return $v, move$

Alpha: the best (highest, at least) value
Beta: the best (lowest, at most) value

Adapted from:
Russell & Norvig, AI: A Modern Approach

Impact of Alpha Beta

- Cuts down on size of game tree searched
 - Best case: $O(b^{(m/2)})$, where b is branching factor and m is depth
 - MINIMAX($O(b^m)$)
- Reduction depends on order of nodes traversed

Game Setting

- Gaming strategies
 - Random
 - Minimax
 - alphabeta
- two automated players
 - First: random
 - Second: alphabeta
 - Code Example: <https://github.com/aimacode/aima-python/blob/master/games4e.ipynb>
- one human, one automated
 - First: human
 - Second: minimax

Heuristic Alpha Beta Tree Search

H-MINIMAX(s) =

- Eval(s, MAX) if IS-CUTOFF(s, d)
- Max a in ACTIONS(S) H-MINIMAX(RESULT(s, a), d+1) if TO-MOVE(s) = MAX
- Min a in ACTIONS(S) H-MINIMAX(RESULT(s, a), d+1) if TO-MOVE(s) = MIN

* Cut-off search at depth d

* Estimate utility of a state to player just like a heuristic function

UTILITY(loss, p) <= EVAL (s, p) <= UTILITY(win, p)

Useful when

- Depth is very high, example Chess
- Good (informative) eval functions exists

Adapted from:
Russell & Norvig, AI: A Modern Approach

Coding Example

- 2-party games code notebook
 - <https://github.com/aimacode/aima-python/blob/master/games4e.ipynb>
 - Has
 - Minimax
 - Alphabeta
 - Heuristic alpha beta

Searching in Larger Games

Game characteristics

- large branching factor – Go (361)
- difficult to get a meaningful evaluation function (heuristics)

Key Idea

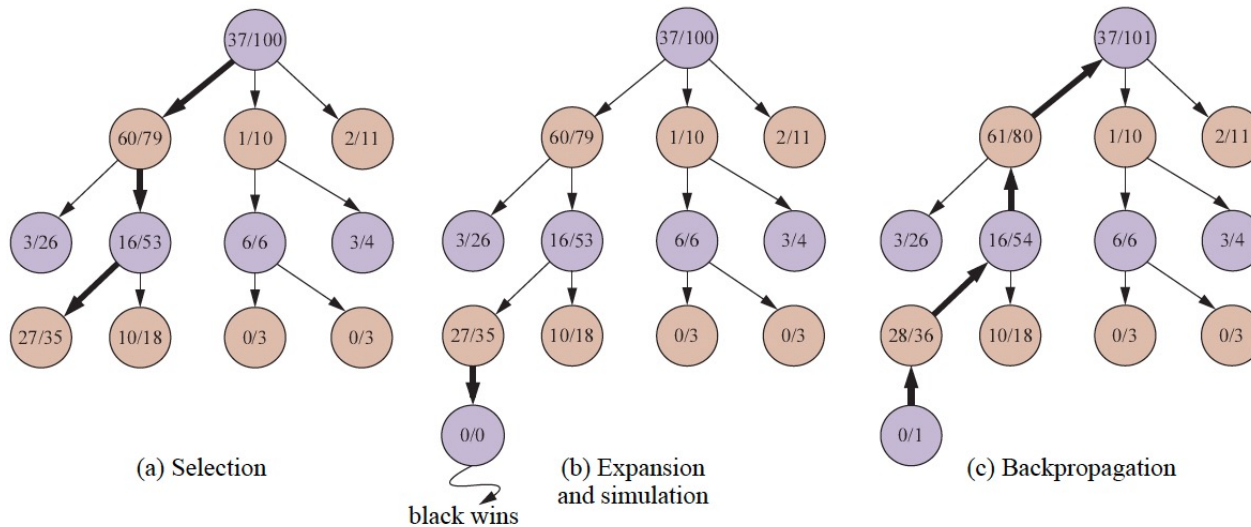
- Value of a state is estimated as the average utility over a number of simulations of complete games from that state
- Get information of good plays by self-play and learning using neural networks

Monte Carlo Tree Search (MCTS)

```
function MONTE-CARLO-TREE-SEARCH(state) returns an action
  tree  $\leftarrow$  NODE(state)
  while IS-TIME-REMAINING() do
    leaf  $\leftarrow$  SELECT(tree)
    child  $\leftarrow$  EXPAND(leaf)
    result  $\leftarrow$  SIMULATE(child)
    BACK-PROPAGATE(result, child)
  return the move in ACTIONS(state) whose node has highest number of playouts
```

Adapted from:
Russell & Norvig, AI: A Modern Approach

Monte Carlo Tree Search (MTCS)



```

function MONTE-CARLO-TREE-SEARCH(state) returns an action
  tree  $\leftarrow$  NODE(state)
  while IS-TIME-REMAINING() do
    leaf  $\leftarrow$  SELECT(tree)
    child  $\leftarrow$  EXPAND(leaf)
    result  $\leftarrow$  SIMULATE(child)
    BACK-PROPAGATE(result, child)
  return the move in ACTIONS(state) whose node has highest number of playouts
  
```

Figure 5.10 One iteration of the process of choosing a move with Monte Carlo tree search (MCTS) using the upper confidence bounds applied to trees (UCT) selection metric, shown after 100 iterations have already been done. In (a) we select moves, all the way down the tree, ending at the leaf node marked 27/35 (for 27 wins for black out of 35 playouts). In (b) we expand the selected node and do a simulation (playout), which ends in a win for black. In (c), the results of the simulation are back-propagated up the tree.

Adapted from:
Russell & Norvig, AI: A Modern Approach

Coding Example

- MCTS code notebook
 - <https://pypi.org/project/mcts-simple/>
 - https://colab.research.google.com/drive/1uYCDn_lymEhexepKfBXcMqiHquyhZpZ5?usp=sharing

Course Project

Project Discussion: What Problem Fascinates You ?

- Data
 - Water
 - Finance
 - ...
- Analytics
 - Search, Optimization, Learning, Planning, ...
- Application
 - Building chatbot
- Users
 - Diverse demographics
 - Diverse abilities
 - Multiple human languages

Project execution in sprints

- Sprint 1: (Sep 12 – Oct 5)
 - **Solving**: Choose a decision problem, identify data, work on solution methods
 - **Human interaction**: Develop a basic chatbot (no AI), no problem focus
- Sprint 2: (Oct 10 – Nov 9)
 - **Solving**: Evaluate your solution on problem
 - **Human interaction**: Integrated your choice of chatbot (rule-based or learning-based) and methods
- Sprint 3: (Nov 14 – 30)
 - **Evaluation**: Comparison of your solver chatbot with an LLM-based alternative, like ChatGPT

Project Discussion: Dates and Deliverables

Project execution in sprints

- Sprint 1: (Sep 12 – Oct 5)
 - **Solving**: Choose a decision problem, identify data, work on solution methods
 - **Human interaction**: Develop a basic chatbot (no AI), no problem focus
- Sprint 2: (Oct 10 – Nov 9)
 - **Solving**: Evaluate your solution on problem
 - **Human interaction**: Integrated your choice of chatbot (rule-based or learning-based) and methods
- Sprint 3: (Nov 14 – 30)
 - **Evaluation**: Comparison of your solver chatbot with an LLM-based alternative, like ChatGPT

- Oct 12, 2023
 - Project checkpoint
 - In-class presentation
- Nov 30, 2023
 - Project report due
- Dec 5 / 7, 2023
 - In-class presentation

Skeleton: A Basic Chatbot

- Run in an infinite loop until the user wants to quit
- Handle any user response
 - User can quit by typing “Quit” or “quit” or just “q”
 - User can enter any other text and the program has to handle it. The program should write back what the user entered and say – “I do not know this information”.
- Handle known user query types // Depends on your project
 - “Tell me about N-queens”, “What is N ?”
 - “Solve for N=4?”
 - “Why is this a solution? ”
- Handle chitchat // Support at least 5, extensible from a file
 - “Hi” => “Hello”
 - ...
- *Store session details in a file*

Illustrative Project

1. **Title:** Solve and explain solving of n-queens puzzle
2. **Key idea:** Show students how a course project will look like
3. **Who will care when done:** students of the course, prospective AI students and teachers
4. **Data need:** n: the size of game; interaction
5. **Methods:** search
6. **Evaluation:** correctness of solution, quality of explanation, appropriateness of chat
7. **Users:** with and without AI background; with and without chess background
8. **Trust issue:** user may not believe in the solution, may find interaction offensive (why queens, not kings? ...)

Project Discussion: Illustration

1. Create a private Github repository called “CSCE58x-Fall2023-<studentname>-Repo”. Share with Instructor (biplav-s) and TA (kausik-l)
2. Create Google folder called “CSCE58x-Fall2023-<studentname>-SharedInfo”. Share with Instructor (prof.biplav@gmail.com) and TA (lakkarajukausik90@gmail.com)
3. Create a Google doc in your Google repo called “Project Plan” and have the following by next class (Sep 5, 2023)

1. **Title:** Solve and explain solving of n-queens puzzle
2. **Key idea:** Show students how a course project will look like
3. **Who will care when done:** students of the course, prospective AI students and teachers
4. **Data need:** n: the size of game; interaction
5. **Methods:** search
6. **Evaluation:** correctness of solution, quality of explanation, appropriateness of chat
7. **Users:** with and without AI background; with and without chess background
8. **Trust issue:** user may not believe in the solution, may find interaction offensive (why queens, not kings? ...)

Project Illustration: N-Queens

- Sprint 1: (Sep 12 – Oct 5)
 - **Solving**: Choose a decision problem, identify data, work on solution methods
 - Method 1: Random solution
 - Method 2: Search – BFS
 - Method 3: Search - ...
 - **Human interaction**: Develop a basic chatbot (no AI) as outlined
 - Deliverable
 - Code structure in Github
 - ./data
 - ./code
 - ./docs
 - ./test
 - Presentation: Make sprint presentation on Oct 12, 2023

Reference: Project Rubric

- **Project results – 60%**
 - Working system ? – 30%
 - Evaluation with results superior to baseline? – 20%
 - Considered related work? – 10%
- **Project efforts – 40%**
 - Project report – 20%
 - Project presentation (updates, final) – 20%
- **Bonus**
 - Challenge level of problem – 10%
 - Instructor discretion – 10%
- **Penalty**
 - Lack of timeliness as per announced policy (right) - up to 30%

Milestones and Penalties

- Oct 12, 2023
 - Project checkpoint
 - In-class presentation
 - **Penalty: presentation not ready by Oct 10, 2023 [-10%]**
- Nov 30, 2023
 - Project report due
 - **Project report not ready by date [-10%]**
- Dec 5 / 7, 2023
 - In-class presentation
 - **Project presentations not ready by Dec 4, 2023 [-10%]**

Lecture 10: Summary

- We talked about
 - Games and Search
 - Minimax
 - Alphabeta
 - Monte Carlo Tree Search

Concluding Section

About Next Lecture – Lecture 11

Lecture 11: Constraints

- Constraints Satisfaction Problems
- Optimization